

Applied AI Engineer with 2+ years of experience building production GenAI systems across agentic workflows, LLM evaluation, RAG, structured outputs, multimodal generation, and AI product automation. Built AI systems for Pixlr’s 6M+ user creative platform, reducing manual intervention by 47%, lowering LLM workflow latency by 33%, and optimizing production inference from 33s to 7.9s p95. Strong in LangGraph, LangChain, MCP-style tool interfaces, FastAPI, PyTorch, Ray Serve, evaluation harnesses.

TECHNICAL EXPERIENCE

Machine Learning & Product Development Engineer Aug 2024 — Jun 2026
Pixlr Pte Ltd | 6M+ user AI photo & design platform Singapore (Remote)

- Designed a production LLM workflow layer for Pixlr’s AI creative tools using LangGraph, structured outputs, prompt versioning, multi-model routing, and moderation fallbacks to support reliable AI design automation at scale.
- Engineered a skill-based agentic workflow using LangGraph, LangChain, and MCP-style tool interfaces to dynamically route creative tasks to specialized skills, reducing manual intervention by 47%.
- Built prompt batching, context reuse, structured response validation, and model-routing logic into the AI workflow harness, reducing median LLM workflow latency by 33% and cost per request by 25%.
- Developed an agent evaluation framework with golden test cases, trace inspection, failure taxonomy, tool-success metrics, and regression checks to measure reliability across multi-step AI workflows.
- Converted ambiguous product ideas into production AI features by collaborating with product managers, engineers, and design stakeholders on API contracts, rollout criteria, validation plans, and user-facing automation behavior.
- Integrated safety and moderation fallbacks into AI generation workflows to detect risky outputs, route uncertain cases to review, and preserve product reliability under latency and cost constraints.
- Optimized underlying multimodal generation services, reducing p95 generation latency from 33s to 7.9s and improving throughput by 120% through serving-layer improvements and GPU workload tuning.

AI Data Expert Dec 2023 — Mar 2024
Scale AI San Francisco (Remote)

- Evaluated** 1,400+ LLM and code-generation outputs using rubric-based criteria across correctness, instruction adherence, reasoning quality, hallucination risk, completeness, and edge-case handling.

Data Science Intern Sep 2023 — Dec 2023
Alpha AI Bengaluru, India

- Performed** exploratory data analysis on customer behavior, product usage, and churn patterns to identify high-risk segments and convert model outputs into actionable retention insights for business users.

EDUCATION

Indian Institute of Information Technology, Allahabad, Bachelor of Technology in Electronics and Communication Engineering
2020 — 2024
CGPA: 8.6 / 10
Merit Scholar for academic excellence
Prayagraj, Uttar Pradesh
2021 — 2022, 2023 — 2024

PROJECTS

SkillMesh: Agentic Workflow Orchestration Platform 2025 – 2026
Tech Stack: LangGraph, LangChain, CrewAI, MCP, Redis

- Architected a skill-based agent orchestration engine using LangGraph contracts to dynamically route tasks, manage shared state, improving task success rate by 42% across a 1.5K+ prompt evaluation suite.
- Implemented a robust shared state management layer using Redis that enabled seamless inter-agent communication, checkpointing, and fault-tolerant recovery, allowing workflows to resume from failure points without loss of context.

Enterprise RAG Control Tower 2024 – 2025
Tech Stack: Python, Vertex AI, FastAPI, FAISS Docker, GCP

- Built a RAG-based analytics layer using Vertex AI, FAISS, FastAPI, and Python to retrieve, rank, and synthesize insights from enterprise documents, reducing P95 query latency by 27% and lowering per-query inference cost by 22%.

SKILLS

Machine Learning	PyTorch, TensorFlow, Scikit-learn, Diffusion Models, SDXL, Multimodal AI, Model Evaluation
Model Serving	FastAPI, Ray Serve, Docker, GPU Inference, Request Batching, Queue Backpressure, Multi-GPU Serving
Evaluation & Reliability	Golden Test Sets, Regression Testing, Trace Inspection, Failure Analysis, Precision/Recall/F1, Latency
GenAI Systems	LLM Workflows, RAG, Structured Outputs, Prompt Versioning, Multi-Model Routing, Moderation Fallbacks
Data & Retrieval	FAISS, Vertex AI Embeddings, Vector Search, Semantic Search, SQL, Pandas, NumPy
Cloud & Systems	AWS, GCP, Redis, CI/CD, Kubernetes, Airflow
Languages	Python, C++, TypeScript, SQL